

Setup batchtools SSH cluster with Boto and Paramiko

I. Make AMI

The AMI is build on Ubuntu 16.04, and has R, sampoll/batchtools and NFS installed.

sampoll/batchtools is different from batchtools only in one line in Worker.R. The worker calls `system.file('/bin/linux-helper', 'batchtools')` on the remote node to find the location of the shell script that executes the actually commands. But in the multiple quoting something is not working. (`/bin/linux-helper` is quoted, the code argument to `Rscript` is quoted, then the entire command is quoted for remote execution.) For this particular case it seemed reasonable to just assume the directory structure is the same on the remote machine and the local (head) node. In general, it is the same, because both nodes are instantiated from the same AMI.

To install from github, we need package `devtools`, and `devtools` needs some special SSL packages.

AMI: ami-1480e86e

```
1 # Build AMI
2 sudo apt-get update
3 sudo apt install r-base-core
4 sudo apt install nfs-kernel-server
5 sudo apt-get install libcurl4-openssl-dev
6 sudo apt-get install libssl-dev # devtools suggests this
7 sudo R
8 > install.packages("devtools", repos="http://lib.stat.cmu.edu/R/CRAN")
9 > install_github("sampoll/batchtools")
```

II. Create cluster with run1.py

The python script `run1.py` creates a security group which allows (1) SSH from anywhere and (2) NFS connection among all group members. Then it creates two EC2 instances and waits for them to come up.

```
1 import boto3
2 from botocore.exceptions import ClientError
3 import pickle
4
5 def create_cluster():
```

```

6
7 """ Cluster configuration is hard-coded into script. Eventually,
8     it should be read from a configuration file. """
9
10 cluster = {
11     'name' : 'cluster000',
12     'vpc' : 'vpc-xxxxxxx',      # usually the user's default VPC
13     'nnode' : 2,
14     'key' : 'XXXXXXX',          # private key is assumed to be in working directory
15     'ami' : 'ami-1480e86e',
16     'type' : 't2.micro',
17 }
18 (sgid,sgn) = create_security_group(cluster)
19 cluster['sgid'] = sgid
20 cluster['sgname'] = sgn
21 ids = create_instances(cluster)
22 cluster['iids'] = ids
23 return cluster
24
25 def save_cluster(cluster):
26
27     """ Pickle cluster for run2.py """
28
29     file = open('cluster.obj', 'wb')
30     pickle.dump(cluster, file)
31     file.close()
32
33 def create_security_group(cluster):
34
35     """ Create a security group for the cluster.
36         Open port 2049 (NFS) for communication within the cluster
37         and port 22 (SSH) to connection from anywhere
38
39         Returns: hash with group name and id"""
40
41     sgn = cluster['name'] + '-sg'
42     client = boto3.client('ec2')
43
44     # If group already exists, delete it. (Note: this will fail if there
45     # is an instance using it.) clusters should be created with
46     # different names.
47
48     sgexists = True
49     try:
50         client.describe_security_groups(GroupNames=[sgn])
51     except ClientError as e:
52         sgexists = False
53     if sgexists:
54         client.delete_security_group(Group_name=sgn)
55
56     security_group_dict = client.create_security_group(

```

```

57     Description='Security Group for SSH Cluster',
58     GroupName=sgn, VpcId=cluster['vpc'])
59
60     sgid = security_group_dict['GroupId']
61     sg = boto3.resource('ec2').SecurityGroup(sgid)
62     p1 = { 'FromPort' : 22, 'ToPort' : 22, 'IpProtocol' : 'tcp' ,
63           'IpRanges' : [ { 'CidrIp' : '0.0.0.0/0', 'Description' : 'Anywhere' } ] }
64     p2 = { 'FromPort' : 2049, 'ToPort' : 2049, 'IpProtocol' : 'tcp' ,
65           'UserIdGroupPairs' : [{ 'GroupId' : sgid } ] }
66     res_auth = sg.authorize_ingress(IpPermissions=[p1, p2])
67     return (sgid, sgn)
68
69 def create_instances(cluster):
70
71     """ Create instances and wait for them to come up.
72     Returns: list of Ids """
73
74     n = cluster['nnode']
75     res = boto3.client('ec2').run_instances(ImageId=cluster['ami'],
76 InstanceType=cluster['type'],
77     KeyName=cluster['key'], MinCount=n, MaxCount=n, SecurityGroupIds=[cluster['sgid']])
78
79     ids = [ inst['InstanceId'] for inst in res['Instances'] ]
80     for id in ids:
81         boto3.resource('ec2').Instance(id).wait_until_running()
82     return ids
83
84 if __name__ == "__main__":
85     cluster = create_cluster()
86     save_cluster(cluster)

```

III. Configure cluster with run2.py

It is possible that run1.py and run2.py could be united into one program. The reason they are separate is that the instance returns true for `wait_until_running()` before the status checks are done. To be safe, I have been waiting until the status checks are complete in the EC2 console.

```

1 # run2.py
2 import boto3
3 from paramiko import client as pclient
4 import sys
5 import pickle
6
7 def init_cluster(cluster):
8
9     """ Get the public and private IPs for each node.
10     Arbitrarily designate one node the head node. """
11
12     filters=[ {'Name': 'instance-id', 'Values': cluster['iids'] } ]

```

```

13 res = boto3.client('ec2').describe_instances(Filters=filters)
14 ips = ipaddresses(res)
15 cluster['head'] = { 'public' : ips[0][1], 'private' : ips[0][0] , 'sg' : ips[0][2] }
16 cluster['compute'] = [ { 'public' : ii[1], 'private' : ii[0] , 'sg' : ii[2] } for ii in
ips[1:] ]
17
18 # extract the IP addresses from the return structure
19 def ipaddresses(rrr):
20     r = []
21     for rr in rrr['Reservations']:
22         instances = rr['Instances']
23         for inst in instances:
24             pvt = inst['PrivateIpAddress']
25             pub = inst['PublicIpAddress']
26             sg = inst['SecurityGroups'][0]['GroupId']
27             r.append( (pvt, pub, sg) )
28     return r
29
30 # NFS export config to be added to /etc/exports on the head node
31 def write_exports(compute):
32     file = open("exports", "w")
33     for c in compute:
34         ss = ("/scratch %(rw, sync, no_root_squash, no_subtree_check)\n" % (c['private']))
35         file.write(ss)
36     file.close()
37
38 # NFS mount config to be added to /etc/fstab on the compute nodes
39 def write_fstab(head):
40     file = open("fstab", "w")
41     ss = ("%s:/scratch /scratch nfs auto, nofail, noatime, nolock, intr, tcp, actimeo=1800 0 0\n" %
head['private'])
42     file.write(ss)
43
44 # batchtools.conf.R for SSH requires IP addresses of compute nodes
45 def write_batchtools_config(compute):
46     file = open("batchtools.conf.R", "w")
47     ss = 'workers = list('
48     for c in compute:
49         ss = ss + 'Worker$new("' + c['private'] + '", ncpus=1)'
50         if c != compute[-1]:
51             ss = ss + ', '
52     ss = ss + ')\n'
53     file.write(ss)
54     file.write('cluster.functions = makeClusterFunctionsSSH(workers)\n')
55     file.close()
56
57 # Execute SSH command remotely
58 def exssh(cl, cmd):
59     print(cmd)
60     (stdin, stdout, stderr) = cl.exec_command(cmd)
61     try:

```

```

62     ex = stdout.channel.recv_exit_status()
63     if ex != 0:
64         print(' !!! Warning: ssh returned non-zero exit status %d ' % (ex))
65 except SSHException:
66     print('SSHException %s' % (cmd))
67
68 def setup_cluster(cluster):
69
70     """ Configure the cluster for use with batchtools SSH option:
71
72         1. Set up passwordless SSH from head node to compute nodes
73         2. Setup up NFS share directory /scratch on all nodes
74         3. Configure /etc/exports on host and /etc/fstab on compute nodes
75            for NFS sharing of /scratch
76         4. Upload batchtools.conf.R to NFS shared directory
77
78     """
79
80     # write IP address-dependent file text in working directory for sftp to cluster nodes
81     write_exports(cluster['compute'])
82     write_fstab(cluster['head'])
83     write_batchtools_config(cluster['compute'])
84
85     # Initiate SSH and SFTP connection to head node
86     headclient = pclient.SSHClient()
87     headclient.set_missing_host_key_policy(pclient.AutoAddPolicy())
88     hip = cluster['head']['public']
89     kfn = cluster['key'] + '.pem'
90     headclient.connect(hip, username='ubuntu', key_filename=kfn)
91     headfclient = headclient.open_sftp()
92
93     # Generate key pair and fetch the public key
94     exssh(headclient, 'sudo rm -f /home/ubuntu/.ssh/id_rsa*')
95     exssh(headclient, 'sudo ssh-keygen -t rsa -f /home/ubuntu/.ssh/id_rsa -q -N ""')
96     exssh(headclient, 'sudo chmod 0400 /home/ubuntu/.ssh/id_rsa')
97     exssh(headclient, 'sudo chown ubuntu:ubuntu /home/ubuntu/.ssh/id_rsa')
98     headfclient.get('/home/ubuntu/.ssh/id_rsa.pub', 'id_rsa.pub')
99
100    # Set up /etc/exports on head node
101    headfclient.put('exports', '/home/ubuntu/exports')
102    exssh(headclient, 'sudo cat /etc/exports /home/ubuntu/exports >/home/ubuntu/tmp')
103    exssh(headclient, 'sudo mv /home/ubuntu/tmp /etc/exports')
104    exssh(headclient, 'sudo rm -f /home/ubuntu/exports')
105    exssh(headclient, 'sudo rm -f /home/ubuntu/tmp')
106
107    # Make /scratch directory on head node
108    exssh(headclient, 'sudo mkdir -p /scratch')
109    exssh(headclient, 'sudo chown nobody:nogroup /scratch')
110    exssh(headclient, 'sudo chmod -R 777 /scratch')
111
112    # Push batchtools.conf.R to /scratch

```

```

113 headfclient.put('batchtools.conf.R', '/scratch/batchtools.conf.R')
114
115 # Restart NFS server on head node
116 exssh(headclient, 'sudo systemctl restart nfs-kernel-server')
117
118 # delete allocated ssh structures
119 headfclient.close()
120 headclient.close()
121
122 for c in cluster['compute']:
123
124     # Initiate SSH and SFTP connection to compute node
125     cip = c['public']
126     nodeclient = pclient.SSHClient()
127     nodeclient.set_missing_host_key_policy(pclient.AutoAddPolicy())
128     nodeclient.connect(cip, username='ubuntu', key_filename=kfn)
129     nodefclient = nodeclient.open_sftp()
130
131     # Propagate public key to compute node
132     nodefclient.put('id_rsa.pub', '/home/ubuntu/.ssh/id_rsa.pub')
133     exssh(nodeclient, 'sudo cat /home/ubuntu/.ssh/id_rsa.pub
>>/home/ubuntu/.ssh/authorized_keys')
134     exssh(nodeclient, 'sudo rm /home/ubuntu/ssh/id_rsa.pub')
135
136     # Make /scratch directory on compute nodes
137     exssh(nodeclient, 'sudo mkdir -p /scratch')
138     exssh(nodeclient, 'sudo chmod -R 777 /scratch')
139
140     # Append mount to /etc/fstab
141     nodefclient.put('fstab', '/home/ubuntu/fstab')
142     exssh(nodeclient, 'sudo cat /etc/fstab /home/ubuntu/fstab >/home/ubuntu/tmp')
143     exssh(nodeclient, 'sudo mv /home/ubuntu/tmp /etc/fstab')
144     exssh(nodeclient, 'sudo rm -f /home/ubuntu/fstab')
145     exssh(nodeclient, 'sudo rm -f /home/ubuntu/tmp')
146
147     # Mount NFS export (Note: /etc/fstab is for automatic mounting on reboot)
148     exssh(nodeclient, 'sudo mount ' + cluster['head']['private'] + ':/scratch /scratch')
149     nodefclient.close()
150     nodeclient.close()
151
152 # Load data from pickle made by run1.py
153 def load_cluster():
154     file = open('cluster.obj', 'rb')
155     cluster = pickle.load(file)
156     file.close()
157     return cluster
158
159 # Pickle cluster for del.py
160 def save_cluster(cluster):
161     file = open('cluster.obj', 'wb')
162     pickle.dump(cluster, file)

```

```

163 file.close()
164
165 if __name__ == '__main__':
166     cluster = load_cluster()
167     init_cluster(cluster)
168     setup_cluster(cluster)
169     save_cluster(cluster)

```

IV. Run tiny example on cluster

ssh to head node

```

1 cd /scratch
2 R
3 R> piApprox = function(n) {
4 +   nums = matrix(runif(2 * n), ncol = 2)
5 +   d = sqrt(nums[,1]^2 + nums[,2]^2)
6 +   4 * mean(d <= 1)
7 + }
8 R> library(batchtools)
9 R> reg <- makeRegistry()      # should indicate registry made with SSH cluster functions
10 R> batchMap(fun = piApprox, n = rep(1e6, 3))
11 R> submitJobs()
12 R> waitForJobs()
13 R> reduceResults(function(x,y) x+y)/3      # 3.1398>

```

V. Delete cluster

Run delete.py

```

1 # delete.py
2 import boto3
3 import pickle
4
5 # load cluster data from pickle
6 file = open('cluster.obj', 'rb')
7 cluster = pickle.load(file)
8 file.close()
9
10 # delete instances and wait for them to terminate
11 client = boto3.client('ec2')
12 res = client.terminate_instances(InstanceIds=cluster['iids'])
13
14 # once instances are terminated, security group can be deleted
15 for id in cluster['iids']:
16     boto3.resource('ec2').Instance(id).wait_until_terminated()
17 client.delete_security_group(Group_name=cluster['sgname'])

```